Automated VM Scaling and Security on GCP

Aryan Kumar M23CSA510

March 1, 2025

GitHub Repository: `https://github.com/Aryank47/gcp-autoscaling`

Abstract

This report outlines step-by-step instructions for creating a VM on Google Cloud Platform (GCP), configuring auto-scaling based on CPU utilization, and implementing security measures such as IAM roles and firewall rules. The architecture ensures scalability, security, and automated testing of the auto-scaling policy.

Contents

1 Introduction

Google Cloud Platform (GCP) is a robust cloud infrastructure for deploying scalable applications. This document demonstrates how to set up a VM instances with auto-scaling and security policies.

2 Step-by-Step Implementation

2.1 Step 1: Create a VM Instance on GCP

2.2 config

```
ZONE=asia-south1-c
MACHINE_TYPE=e2-micro
TEMPLATE_NAME=test-template
MIG_NAME=test-mig
TARGET_CPU=0.7
MAX_INSTANCES=3
LOAD_DURATION=300
```

2.3 Authenticate and Initialize the CLI

```
gcloud auth login
gcloud config set project [PROJECT_ID]
gcloud config set compute/zone [ZONE]
```

2.4 Create a VM Instance

```
gcloud compute instance-templates create $TEMPLATE_NAME \
        --machine-type=$MACHINE_TYPE \
        --image-family=debian-11 \
        --image-project=debian-cloud \
        --tags=allow-ssh,allow-http \
        --metadata=startup-script='
#!/bin/bash
apt-update
apt-install-y-stress-ng-python3
(
cat-<<-EOF
import-os

def-generate_stress(duration=300):
----command-=-f"stress-ng-----cpu-2-----timeout-{duration}s"
----os.system(command)

if-__name__-==-"__main__":
----generate_stress()
EOF
)->-/usr/local/bin/load_generator.py
chmod-+x-/usr/local/bin/load_generator.py
--------'
```

2.5 Verify the Instance

```
gcloud compute instances list
```

2.6 Step 2: Configure Auto-Scaling Policy

2.7 Create a Managed Instance Group (MIG)

```
gcloud compute instance-groups managed create $MIG_NAME \
        --template=$TEMPLATE_NAME \
        --base-instance-name=test-vm \
        --zone=$ZONE \
        --size=1
```

2.8 Set Auto-Scaling Policy

```
gcloud compute instance-groups managed set-autoscaling $MIG_NAME \
        --max-num-replicas=$MAX_INSTANCES \
        --min-num-replicas=1 \
        --target-cpu-utilization=$TARGET_CPU \
        --zone=$ZONE
```

2.9 Verify Auto-Scaling Configuration

```
gcloud compute instance-groups managed describe my-mig --zone=[ZONE]
```

2.10 Step 3: Implement Security Measures

2.11 Create Firewall Rules

```
gcloud compute firewall-rules create allow-ssh \
  --allow tcp:22 \
  --target-tags=allow-ssh

gcloud compute firewall-rules create allow-http \
  --allow tcp:80 \
  --target-tags=allow-http

gcloud compute firewall-rules create deny-all \
  --deny all \
  --priority=65534
```

2.12 Configure IAM Roles

```
gcloud iam service-accounts create compute-reader --display-name="Compute Reader"
gcloud projects add-iam-policy-binding [PROJECT_ID] \
  --member="serviceAccount:compute-reader@[PROJECT_ID].iam.gserviceaccount.com" \
  --role="roles/compute.instanceViewer"
```

2.13 Verify IAM Policy

```
gcloud projects get-iam-policy [PROJECT_ID] --flatten="bindings[].members"
```

2.14 Step 4: Test Auto-Scaling

2.15 Generate CPU Load

```
gcloud compute ssh my-vm --zone=[ZONE] --command="stress-ng --cpu 2 --timeout 300s"
```

2.16 Monitor Scaling Activity

```
gcloud compute instance-groups managed list-instances my-mig --zone=[ZONE]
```

3 Architecture Diagram

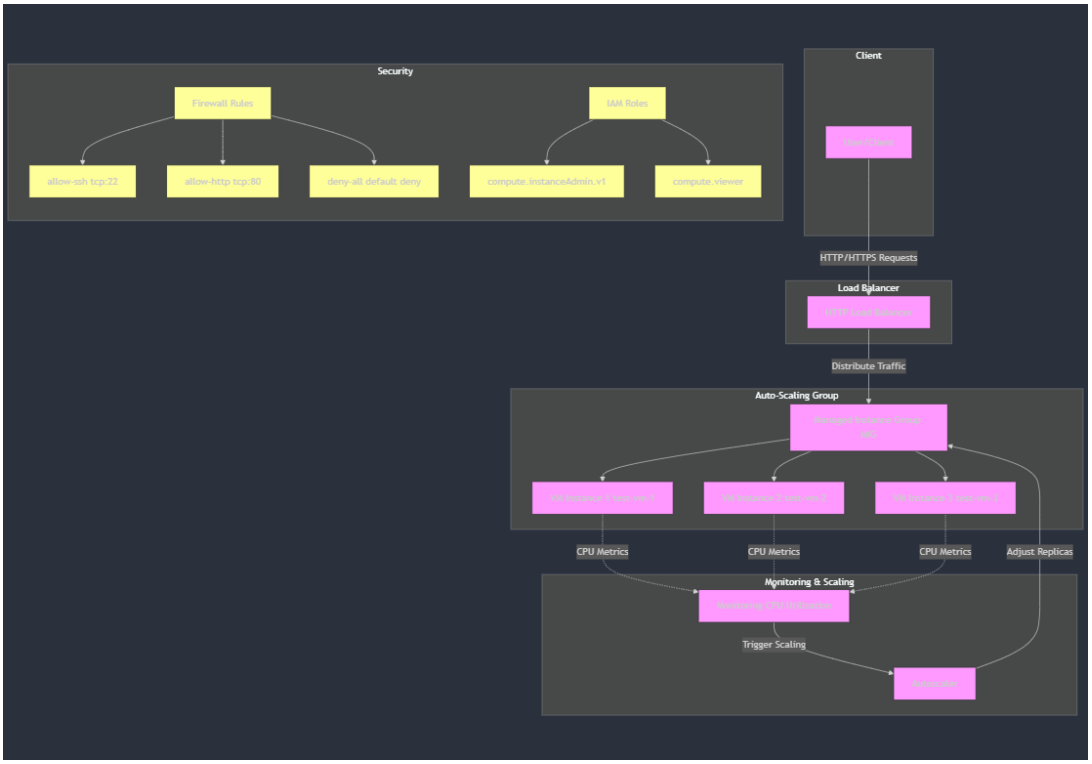

Figure 1: **GCP Auto-Scaling Architecture**

4 CPU Utilization Metrics

4.1 Initial Load Generation

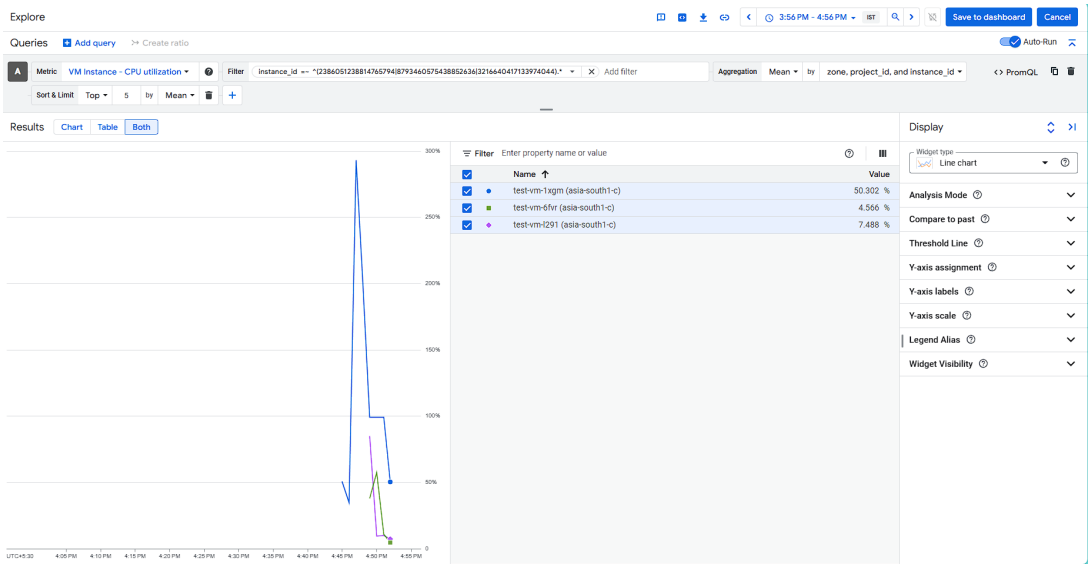

Figure 2: Initial CPU utilization spike during load generation

Description: This graph shows the CPU utilization of the VM instances during the initial load generation phase. The blue line represents the primary instance experiencing a significant spike in CPU usage, triggering the auto-scaling policy.

4.2 Auto-Scaling Response

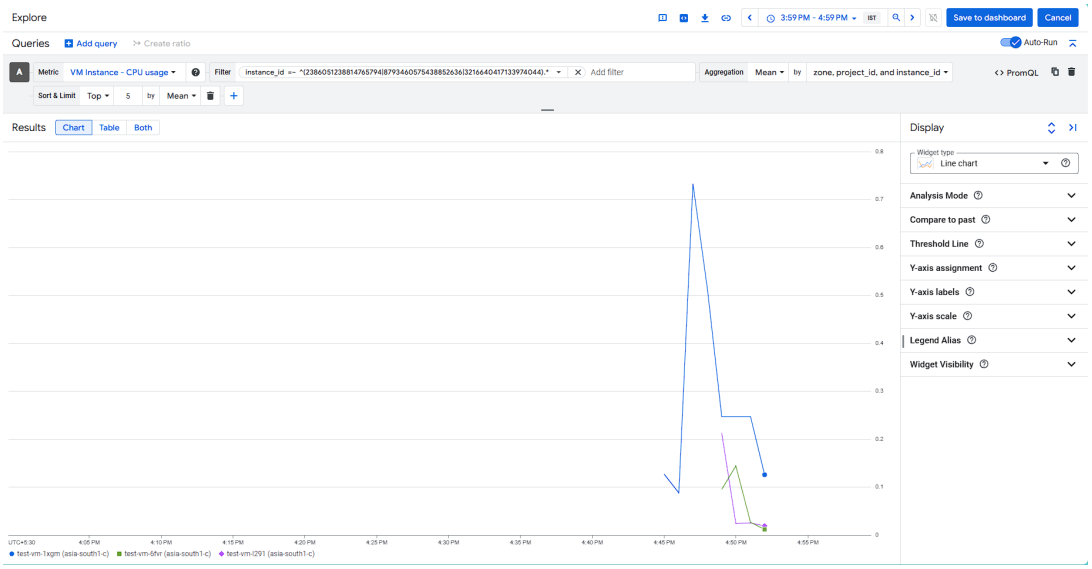

Figure 3: Auto-scaling response and load distribution

Description: This graph illustrates the auto-scaling response where additional instances are created to handle the load. The different colored lines represent multiple VM instances sharing the load, demonstrating the effectiveness of the auto-scaling configuration.

5 Terminal Log Screenshots

5.1 Auto-Scaling Test Execution

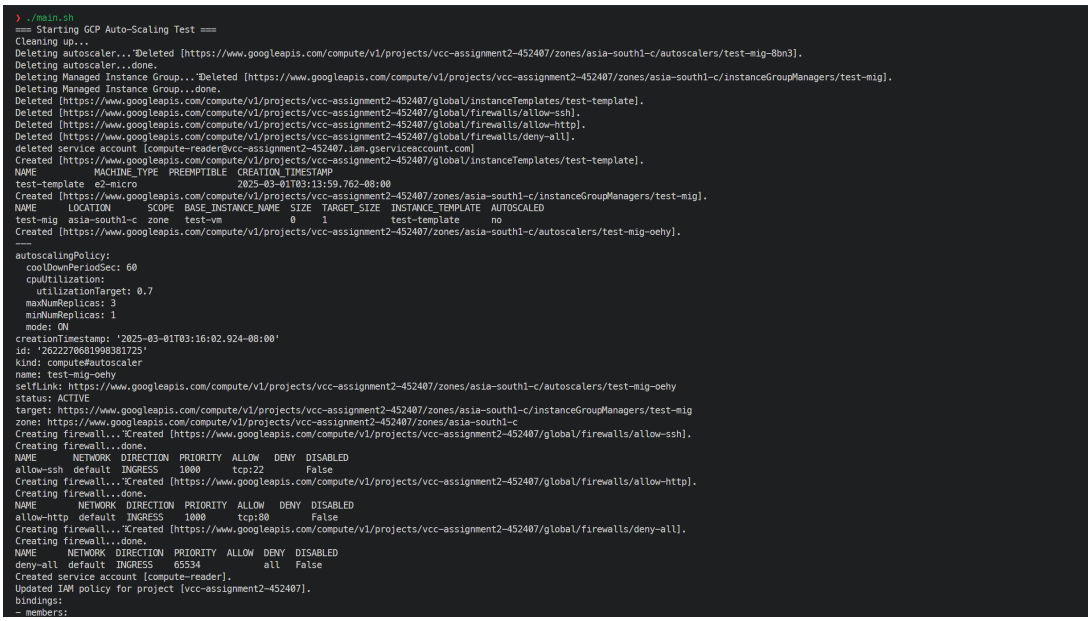

Figure 4: Auto-scaling test execution log showing instance scaling from 1 to 3

Description: This screenshot shows the terminal output during the auto-scaling test. The script starts with 1 instance, generates CPU load using `stress-ng`, and scales up to 3 instances as the load increases. The firewall rules and IAM policies are verified at the end.

5.2 Resource Cleanup

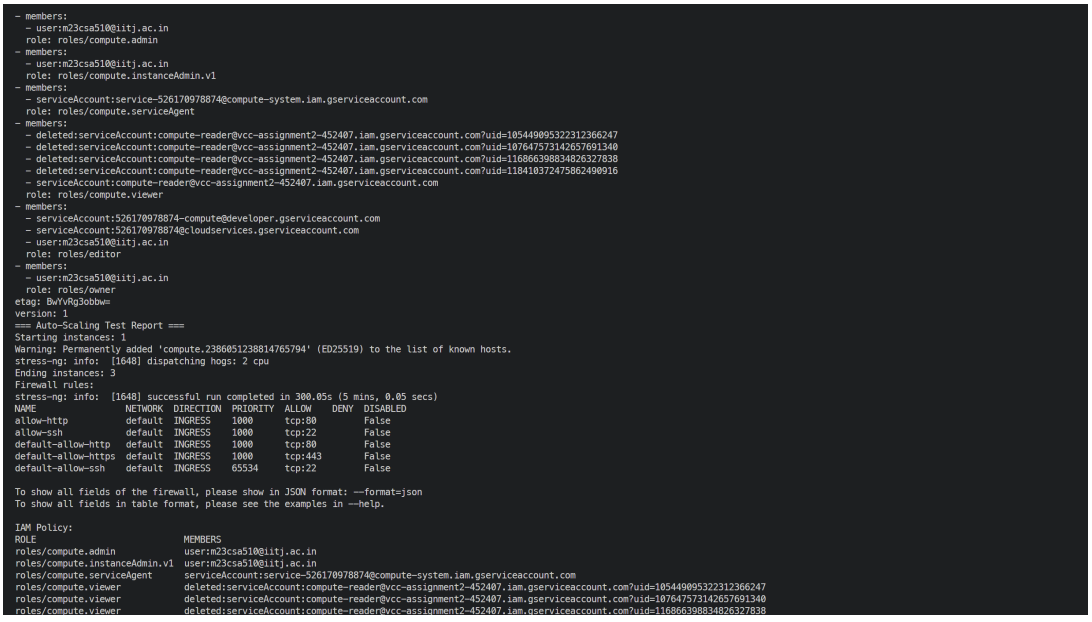

Figure 5: Resource cleanup log showing deletion of GCP resources

Description: This screenshot shows the terminal output during resource cleanup. All created resources including the autoscaler, managed instance group, instance template, firewall rules, and service account are deleted to avoid billing charges.

5.3 IAM Policy Verification

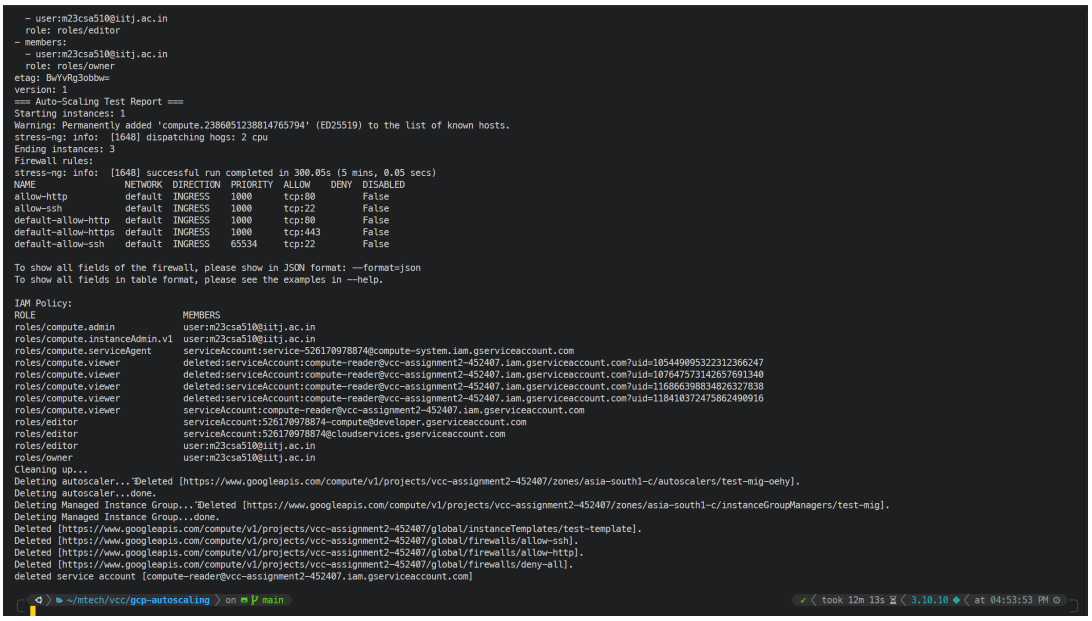

Figure 6: IAM policy verification log showing assigned roles and members

Description: This screenshot shows the output of the IAM policy verification command. It lists all the roles assigned to different members, including the compute admin, instance admin, service agent, and viewer roles for the service account.

6 Best Practices

- Use the default VPC network for simplicity.
- Keep zones consistent across resources.
- Delete resources after testing to avoid costs.

7 Conclusion

This implementation provides a scalable and secure VM setup on GCP. The auto-scaling policy ensures resource optimization, while IAM and firewall rules enhance security. Always validate your setup by triggering load and monitoring the environment.